

```

# Install libraries if needed
!pip install pandas numpy scikit-learn matplotlib seaborn

# Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix

Requirement already satisfied: pandas in
/usr/local/lib/python3.11/dist-packages (2.2.2)
Requirement already satisfied: numpy in
/usr/local/lib/python3.11/dist-packages (1.26.4)
Requirement already satisfied: scikit-learn in
/usr/local/lib/python3.11/dist-packages (1.6.1)
Requirement already satisfied: matplotlib in
/usr/local/lib/python3.11/dist-packages (3.10.0)
Requirement already satisfied: seaborn in
/usr/local/lib/python3.11/dist-packages (0.13.2)
Requirement already satisfied: python-dateutil<=2.8.2 in
/usr/local/lib/python3.11/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz<=2020.1 in
/usr/local/lib/python3.11/dist-packages (from pandas) (2024.2)
Requirement already satisfied: tzdata<=2022.7 in
/usr/local/lib/python3.11/dist-packages (from pandas) (2025.1)
Requirement already satisfied: scipy<=1.6.0 in
/usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.13.1)
Requirement already satisfied: joblib<=1.2.0 in
/usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl<=3.1.0 in
/usr/local/lib/python3.11/dist-packages (from scikit-learn) (3.5.0)
Requirement already satisfied: contourpy<=1.0.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3.1)
Requirement already satisfied: cycler<=0.10 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools<=4.22.0 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (4.55.7)
Requirement already satisfied: kiwisolver<=1.3.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (1.4.8)
Requirement already satisfied: packaging<=20.0 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (24.2)
Requirement already satisfied: pillow<=8 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (11.1.0)
Requirement already satisfied: pyparsing<=2.3.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (3.2.1)
Requirement already satisfied: six<=1.5 in

```

```
/usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2-  
>pandas) (1.17.0)
```

```
df=pd.read_csv('traffic_accidents.csv')  
print(df)
```

	weather_condition \	crash_date	traffic_control_device
0	CLEAR	07/29/2023 01:00:00 PM	TRAFFIC SIGNAL
1	CLEAR	08/13/2023 12:11:00 AM	TRAFFIC SIGNAL
2	CLEAR	12/09/2021 10:30:00 AM	TRAFFIC SIGNAL
3	CLEAR	08/09/2023 07:55:00 PM	TRAFFIC SIGNAL
4	CLEAR	08/19/2023 02:55:00 PM	TRAFFIC SIGNAL
...	
120257	CLEAR	01/11/2023 08:12:00 AM	TRAFFIC SIGNAL
120258	CLEAR	07/16/2018 07:58:00 AM	NO CONTROLS
120259	CLEAR	10/17/2023 03:33:00 PM	TRAFFIC SIGNAL
120260	RAIN	10/19/2020 08:47:00 PM	TRAFFIC SIGNAL
120261	CLEAR	09/28/2019 04:50:00 PM	TRAFFIC SIGNAL

	lighting_condition	first_crash_type \
0	DAYLIGHT	TURNING
1	DARKNESS, LIGHTED ROAD	TURNING
2	DAYLIGHT	REAR END
3	DAYLIGHT	ANGLE
4	DAYLIGHT	REAR END
...
120257	DAYLIGHT	TURNING
120258	DAYLIGHT	TURNING
120259	DAYLIGHT	TURNING
120260	DARKNESS, LIGHTED ROAD	REAR END
120261	DAYLIGHT	TURNING

	trafficway_type	alignment \
0	NOT DIVIDED	STRAIGHT AND LEVEL
1	FOUR WAY	STRAIGHT AND LEVEL
2	T-INTERSECTION	STRAIGHT AND LEVEL
3	FOUR WAY	STRAIGHT AND LEVEL
4	T-INTERSECTION	STRAIGHT AND LEVEL

...
120257	FOUR WAY	STRAIGHT AND LEVEL
120258	DIVIDED - W/MEDIAN (NOT RAISED)	STRAIGHT AND LEVEL
120259	NOT DIVIDED	STRAIGHT AND LEVEL
120260	NOT DIVIDED	STRAIGHT AND LEVEL
120261	DIVIDED - W/MEDIAN BARRIER	STRAIGHT AND LEVEL

	roadway_surface_cond	road_defect	
crash_type \			
0	UNKNOWN	UNKNOWN	NO INJURY / DRIVE AWAY
1	DRY	NO DEFECTS	NO INJURY / DRIVE AWAY
2	DRY	NO DEFECTS	NO INJURY / DRIVE AWAY
3	DRY	NO DEFECTS	INJURY AND / OR TOW DUE TO CRASH
4	UNKNOWN	UNKNOWN	NO INJURY / DRIVE AWAY
...

...			
120257	DRY	NO DEFECTS	INJURY AND / OR TOW DUE TO CRASH
120258	DRY	NO DEFECTS	INJURY AND / OR TOW DUE TO CRASH
120259	DRY	NO DEFECTS	NO INJURY / DRIVE AWAY
120260	WET	UNKNOWN	NO INJURY / DRIVE AWAY
120261	DRY	NO DEFECTS	INJURY AND / OR TOW DUE TO CRASH

	...	most_severe_injury	injuries_total	injuries_fatal \
0	...	NO INDICATION OF INJURY	0.0	0.0
1	...	NO INDICATION OF INJURY	0.0	0.0
2	...	NO INDICATION OF INJURY	0.0	0.0
3	...	NONINCAPACITATING INJURY	5.0	0.0
4	...	NO INDICATION OF INJURY	0.0	0.0
...
120257	...	NONINCAPACITATING INJURY	1.0	0.0
120258	...	NO INDICATION OF INJURY	0.0	0.0
120259	...	NO INDICATION OF INJURY	0.0	0.0
120260	...	NO INDICATION OF INJURY	0.0	0.0
120261	...	NO INDICA	NaN	NaN

	injuries_incapacitating	injuries_non_incapacitating \
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	5.0

4	0.0	0.0
...
120257	0.0	1.0
120258	0.0	0.0
120259	0.0	0.0
120260	0.0	0.0
120261	NaN	NaN

	injuries_reported_not_evident	injuries_no_indication
crash_hour \		
0	0.0	3.0
13.0		
1	0.0	2.0
0.0		
2	0.0	3.0
10.0		
3	0.0	0.0
19.0		
4	0.0	3.0
14.0		
...
...		
120257	0.0	2.0
8.0		
120258	0.0	2.0
7.0		
120259	0.0	2.0
15.0		
120260	0.0	3.0
20.0		
120261	NaN	NaN
NaN		

	crash_day_of_week	crash_month
0	7.0	7.0
1	1.0	8.0
2	5.0	12.0
3	4.0	8.0
4	7.0	8.0
...
120257	4.0	1.0
120258	2.0	7.0
120259	3.0	10.0
120260	2.0	10.0
120261	NaN	NaN

[120262 rows x 24 columns]

`print(df.columns)`

```
Index(['crash_date', 'traffic_control_device', 'weather_condition',
      'lighting_condition', 'first_crash_type', 'trafficway_type',
      'alignment', 'roadway_surface_cond', 'road_defect',
      'crash_type',
      'intersection_related_i', 'damage', 'prim_contributory_cause',
      'num_units', 'most_severe_injury', 'injuries_total',
      'injuries_fatal',
      'injuries_incapacitating', 'injuries_non_incapacitating',
      'injuries_reported_not_evident', 'injuries_no_indication',
      'crash_hour',
      'crash_day_of_week', 'crash_month'],
      dtype='object')
```

```
# Check for missing values
```

```
print(df.isnull().sum())
```

```
# Fill or drop missing values (example: filling with mode for
categorical columns)
```

```
for col in df.select_dtypes(include='object').columns:
    df[col].fillna(df[col].mode()[0], inplace=True)
```

```
# Fill numeric columns with median
```

```
for col in df.select_dtypes(include='number').columns:
    df[col].fillna(df[col].median(), inplace=True)
```

```
# Encode categorical variables
```

```
label_encoders = {}
```

```
for col in df.select_dtypes(include='object').columns:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le
```

```
# Display cleaned data
```

```
df.info()
```

```
crash_date                0
traffic_control_device    0
weather_condition         0
lighting_condition        0
first_crash_type          0
trafficway_type            0
alignment                 0
roadway_surface_cond      0
road_defect               0
crash_type                0
intersection_related_i    0
damage                    0
prim_contributory_cause   0
num_units                 0
most_severe_injury        0
```

```

injuries_total          1
injuries_fatal          1
injuries_incapacitating 1
injuries_non_incapacitating 1
injuries_reported_not_evident 1
injuries_no_indication  1
crash_hour              1
crash_day_of_week       1
crash_month             1
dtype: int64

```

<ipython-input-10-61791cad35b5>:6: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df[col].fillna(df[col].mode()[0], inplace=True)
```

<ipython-input-10-61791cad35b5>:10: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df[col].fillna(df[col].median(), inplace=True)
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 120262 entries, 0 to 120261
Data columns (total 24 columns):

```

#	Column	Non-Null Count	Dtype
0	crash_date	120262 non-null	int64
1	traffic_control_device	120262 non-null	int64
2	weather_condition	120262 non-null	int64
3	lighting_condition	120262 non-null	int64
4	first_crash_type	120262 non-null	int64

5	trafficway_type	120262	non-null	int64
6	alignment	120262	non-null	int64
7	roadway_surface_cond	120262	non-null	int64
8	road_defect	120262	non-null	int64
9	crash_type	120262	non-null	int64
10	intersection_related_i	120262	non-null	int64
11	damage	120262	non-null	int64
12	prim_contributory_cause	120262	non-null	int64
13	num_units	120262	non-null	int64
14	most_severe_injury	120262	non-null	int64
15	injuries_total	120262	non-null	float64
16	injuries_fatal	120262	non-null	float64
17	injuries_incapacitating	120262	non-null	float64
18	injuries_non_incapacitating	120262	non-null	float64
19	injuries_reported_not_evident	120262	non-null	float64
20	injuries_no_indication	120262	non-null	float64
21	crash_hour	120262	non-null	float64
22	crash_day_of_week	120262	non-null	float64
23	crash_month	120262	non-null	float64

dtypes: float64(9), int64(15)

memory usage: 22.0 MB

Define features and target

X = df.drop(['crash_type'], axis=1) *# Replace 'crash_type' if using another target*

y = df['crash_type']

Split data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

Initialize and train model

model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X_train, y_train)

Make predictions

y_pred = model.predict(X_test)

Classification report


print(classification_report(y_test, y_pred))

Confusion matrix

conf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.87	0.75	0.81	10543
1	0.82	0.91	0.87	13510
accuracy			0.84	24053
macro avg	0.85	0.83	0.84	24053
weighted avg	0.84	0.84	0.84	24053



```

from sklearn.model_selection import GridSearchCV

# Define parameter grid for tuning
param_grid = {
    'n_estimators': [50, 100, 150],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Initialize GridSearchCV
grid_search =
GridSearchCV(estimator=RandomForestClassifier(random_state=42),
              param_grid=param_grid,
              cv=3,

```



```

n_jobs=-1,
verbose=2)

# Fit model to training data
grid_search.fit(X_train, y_train)

# Display best parameters
print(f"Best Parameters: {grid_search.best_params_}")

Fitting 3 folds for each of 108 candidates, totalling 324 fits
Best Parameters: {'max_depth': None, 'min_samples_leaf': 4,
'min_samples_split': 2, 'n_estimators': 100}

# Use the best estimator found by GridSearchCV
best_model = grid_search.best_estimator_

# Predict using the tuned model
y_pred_tuned = best_model.predict(X_test)

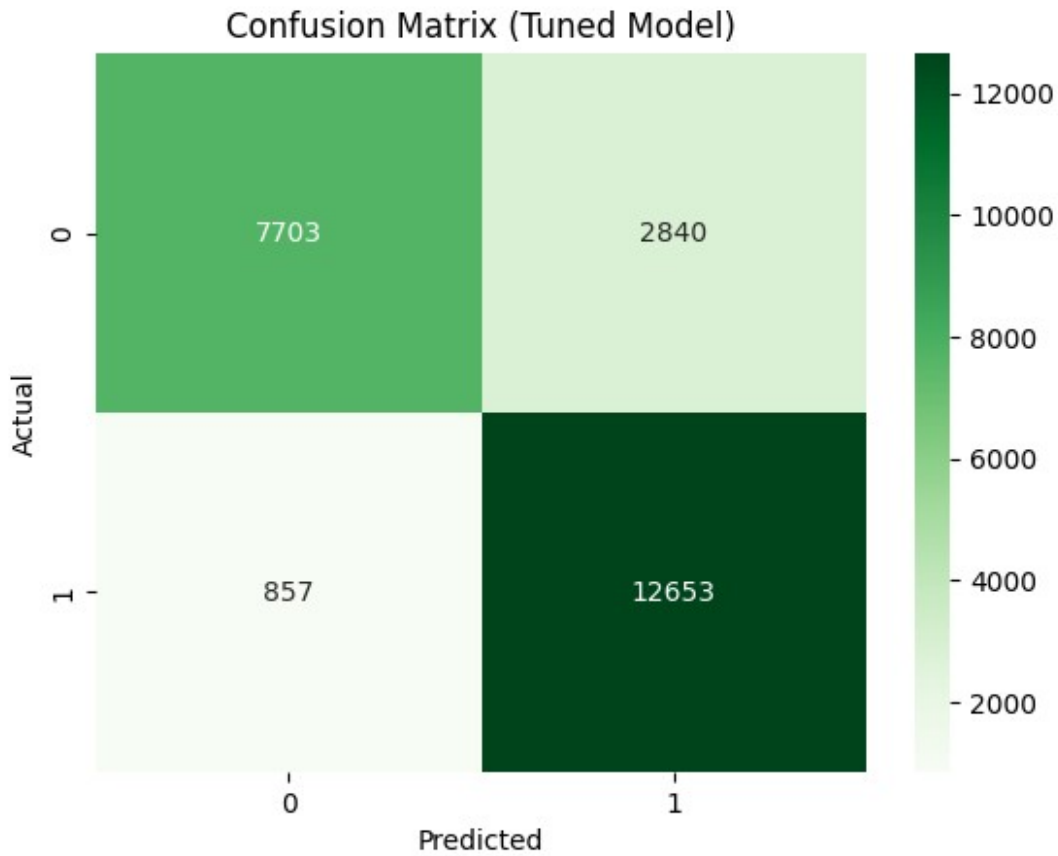
# Evaluate performance
from sklearn.metrics import classification_report, confusion_matrix

print(classification_report(y_test, y_pred_tuned))

# Confusion matrix
conf_matrix_tuned = confusion_matrix(y_test, y_pred_tuned)
sns.heatmap(conf_matrix_tuned, annot=True, fmt='d', cmap='Greens')
plt.title('Confusion Matrix (Tuned Model)')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

```

	precision	recall	f1-score	support
0	0.90	0.73	0.81	10543
1	0.82	0.94	0.87	13510
accuracy			0.85	24053
macro avg	0.86	0.83	0.84	24053
weighted avg	0.85	0.85	0.84	24053



```
# Get feature importances from the model
importances = best_model.feature_importances_
features = X.columns

# Create DataFrame for visualization
feature_importances_df = pd.DataFrame({'Feature': features,
                                       'Importance': importances})
feature_importances_df =
feature_importances_df.sort_values(by='Importance', ascending=False)

# Display top features
print(feature_importances_df.head(10))

# Plot feature importances
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature',
            data=feature_importances_df.head(10), palette='viridis')
plt.title('Top 10 Feature Importances')
plt.show()
```


	Feature	Importance
14	injuries_total	0.217534
13	most_severe_injury	0.195029

17	injuries_non_incapacitating	0.084180
4	first_crash_type	0.078058
19	injuries_no_indication	0.070962
10	damage	0.061369
18	injuries_reported_not_evident	0.047949
11	prim_contributory_cause	0.043100
0	crash_date	0.040668
20	crash_hour	0.033366

<ipython-input-18-67c2ae666497>:14: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Importance', y='Feature',
data=feature_importances_df.head(10), palette='viridis')
```



```
import joblib

# Save the model
joblib.dump(best_model, 'accident_prediction_model.pkl')

['accident_prediction_model.pkl']

# Load model
model = joblib.load('accident_prediction_model.pkl')

# Get the feature names used during training
training_features = model.feature_names_in_
```

```

# Example input (replace with user input from chatbot)
# Ensure all features used during training are included and in the
correct order
user_input = pd.DataFrame({
    feature: [0] for feature in training_features # Initialize all
features to 0
})

# Update user input with actual values (replace with user input from
chatbot)
user_input['weather_condition'] = [1] # Example encoded value
user_input['lighting_condition'] = [2]
user_input['roadway_surface_cond'] = [0]
user_input['crash_hour'] = [14]
user_input['traffic_control_device'] = [3]
# ... Add other necessary features here ...

```

```

# Predict
prediction = model.predict(user_input)
print(f'Predicted Crash Type: {prediction[0]}')

```

Predicted Crash Type: 1

pip install streamlit

Collecting streamlit

```

Downloading streamlit-1.41.1-py2.py3-none-any.whl.metadata (8.5 kB)
Requirement already satisfied: altair<6,>=4.0 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (5.5.0)
Requirement already satisfied: blinker<2,>=1.0.0 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (1.9.0)
Requirement already satisfied: cachetools<6,>=4.0 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (5.5.1)
Requirement already satisfied: click<9,>=7.0 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (8.1.8)
Requirement already satisfied: numpy<3,>=1.23 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (1.26.4)
Requirement already satisfied: packaging<25,>=20 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (24.2)
Requirement already satisfied: pandas<3,>=1.4.0 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (2.2.2)
Requirement already satisfied: pillow<12,>=7.1.0 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (11.1.0)
Requirement already satisfied: protobuf<6,>=3.20 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (4.25.6)
Requirement already satisfied: pyarrow>=7.0 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (17.0.0)
Requirement already satisfied: requests<3,>=2.27 in

```

```
/usr/local/lib/python3.11/dist-packages (from streamlit) (2.32.3)
Requirement already satisfied: rich<14,>=10.14.0 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (13.9.4)
Requirement already satisfied: tenacity<10,>=8.1.0 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (9.0.0)
Requirement already satisfied: toml<2,>=0.10.1 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (0.10.2)
Requirement already satisfied: typing-extensions<5,>=4.3.0 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (4.12.2)
Collecting watchdog<7,>=2.1.5 (from streamlit)
  Downloading watchdog-6.0.0-py3-none-
manylinux2014_x86_64.whl.metadata (44 kB)
----- 44.3/44.3 kB 2.4 MB/s eta
0:00:00
Requirement already satisfied: gitpython!=3.1.19,<4,>=3.0.7 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (3.1.44)
Collecting pydeck<1,>=0.8.0b4 (from streamlit)
  Downloading pydeck-0.9.1-py2.py3-none-any.whl.metadata (4.1 kB)
Requirement already satisfied: tornado<7,>=6.0.3 in
/usr/local/lib/python3.11/dist-packages (from streamlit) (6.4.2)
Requirement already satisfied: jinja2 in
/usr/local/lib/python3.11/dist-packages (from altair<6,>=4.0-
>streamlit) (3.1.5)
Requirement already satisfied: jsonschema>=3.0 in
/usr/local/lib/python3.11/dist-packages (from altair<6,>=4.0-
>streamlit) (4.23.0)
Requirement already satisfied: narwhals>=1.14.2 in
/usr/local/lib/python3.11/dist-packages (from altair<6,>=4.0-
>streamlit) (1.24.1)
Requirement already satisfied: gitdb<5,>=4.0.1 in
/usr/local/lib/python3.11/dist-packages (from gitpython!
=3.1.19,<4,>=3.0.7->streamlit) (4.0.12)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.11/dist-packages (from pandas<3,>=1.4.0-
>streamlit) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.11/dist-packages (from pandas<3,>=1.4.0-
>streamlit) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in
/usr/local/lib/python3.11/dist-packages (from pandas<3,>=1.4.0-
>streamlit) (2025.1)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.11/dist-packages (from requests<3,>=2.27-
>streamlit) (3.4.1)
Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.11/dist-packages (from requests<3,>=2.27-
>streamlit) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.11/dist-packages (from requests<3,>=2.27-
```

```

>streamlit) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.11/dist-packages (from requests<3,>=2.27-
>streamlit) (2024.12.14)
Requirement already satisfied: markdown-it-py>=2.2.0 in
/usr/local/lib/python3.11/dist-packages (from rich<14,>=10.14.0-
>streamlit) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in
/usr/local/lib/python3.11/dist-packages (from rich<14,>=10.14.0-
>streamlit) (2.18.0)
Requirement already satisfied: smmap<6,>=3.0.1 in
/usr/local/lib/python3.11/dist-packages (from gitdb<5,>=4.0.1-
>gitpython!=3.1.19,<4,>=3.0.7->streamlit) (5.0.2)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.11/dist-packages (from jinja2->altair<6,>=4.0-
>streamlit) (3.0.2)
Requirement already satisfied: attrs>=22.2.0 in
/usr/local/lib/python3.11/dist-packages (from jsonschema>=3.0-
>altair<6,>=4.0->streamlit) (25.1.0)
Requirement already satisfied: jsonschema-specifications>=2023.03.6 in
/usr/local/lib/python3.11/dist-packages (from jsonschema>=3.0-
>altair<6,>=4.0->streamlit) (2024.10.1)
Requirement already satisfied: referencing>=0.28.4 in
/usr/local/lib/python3.11/dist-packages (from jsonschema>=3.0-
>altair<6,>=4.0->streamlit) (0.36.2)
Requirement already satisfied: rpds-py>=0.7.1 in
/usr/local/lib/python3.11/dist-packages (from jsonschema>=3.0-
>altair<6,>=4.0->streamlit) (0.22.3)
Requirement already satisfied: mdurl~=0.1 in
/usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0-
>rich<14,>=10.14.0->streamlit) (0.1.2)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2-
>pandas<3,>=1.4.0->streamlit) (1.17.0)
Downloading streamlit-1.41.1-py2.py3-none-any.whl (9.1 MB)
0:00:00 9.1/9.1 MB 51.8 MB/s eta
0:00:00 6.9/6.9 MB 45.8 MB/s eta
anylinux2014_x86_64.whl (79 kB)
0:00:00 79.1/79.1 kB 5.0 MB/s eta
lit
Successfully installed pydeck-0.9.1 streamlit-1.41.1 watchdog-6.0.0

# List of features used during training
print(list(X.columns))

['crash_date', 'traffic_control_device', 'weather_condition',
'lighting_condition', 'first_crash_type', 'trafficway_type',

```

```
'alignment', 'roadway_surface_cond', 'road_defect',  
'intersection_related_i', 'damage', 'prim_contributory_cause',  
'num_units', 'most_severe_injury', 'injuries_total', 'injuries_fatal',  
'injuries_incapacitating', 'injuries_non_incapacitating',  
'injuries_reported_not_evident', 'injuries_no_indication',  
'crash_hour', 'crash_day_of_week', 'crash_month']
```

```
# Before training the model
```

```
X = X.drop(['crash_date'], axis=1)
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
model = RandomForestClassifier()
```

```
model.fit(X, y)
```

```
# Save the retrained model
```

```
import joblib
```

```
joblib.dump(model, 'accident_prediction_model.pkl')
```

```
['accident_prediction_model.pkl']
```

```
# Load the saved model
```

```
import joblib
```

```
import pandas as pd
```

```
model = joblib.load('accident_prediction_model.pkl')
```

```
# Simple chatbot function
```

```
def accident_prediction_chatbot():
```

```
    print("Welcome to the Accident Prediction Chatbot!")
```

```
    print("Please provide the following information:\n")
```

```
    # User inputs
```

```
    weather_condition = int(input("Weather Condition (0=Clear, 1=Rain, 2=Snow, etc.): "))
```

```
    lighting_condition = int(input("Lighting Condition (0=Daylight, 1=Dark, etc.): "))
```

```
    roadway_surface_cond = int(input("Road Surface Condition (0=Dry, 1=Wet, etc.): "))
```

```
    crash_hour = int(input("Crash Hour (0-23): "))
```

```
    traffic_control_device = int(input("Traffic Control Device (0=None, 1=Stop Sign, etc.): "))
```

```
    first_crash_type = int(input("First Crash Type (0=Rear-end, 1=Angle, etc.): "))
```

```
    alignment = int(input("Alignment (0=Straight, 1=Curve, etc.): "))
```

```
    prim_contributory_cause = int(input("Primary Cause (0=Speeding, 1=Distracted Driving, etc.): "))
```

```
    most_severe_injury = int(input("Most Severe Injury (0=None, 1=Minor, 2=Severe, 3=Fatal): "))
```

```
# Default values for missing features
```

```

crash_date = '2025-01-01' # Placeholder date
crash_day_of_week = 4      # Assuming Thursday (adjust as needed)
crash_month = 1            # January
damage = 0                 # No damage (adjust based on dataset
encoding)
num_units = 1              # Assuming a single vehicle
intersection_related_i = 0 # Not at intersection
injuries_total = 0
injuries_fatal = 0
injuries_incapacitating = 0
injuries_non_incapacitating = 0
injuries_reported_not_evident = 0
injuries_no_indication = 0
road_defect = 0            # Assuming no defect
trafficway_type = 0        # Default trafficway type

# Create DataFrame with all required features
user_data = pd.DataFrame({
    'weather_condition': [weather_condition],
    'lighting_condition': [lighting_condition],
    'roadway_surface_cond': [roadway_surface_cond],
    'crash_hour': [crash_hour],
    'traffic_control_device': [traffic_control_device],
    'first_crash_type': [first_crash_type],
    'alignment': [alignment],
    'prim_contributory_cause': [prim_contributory_cause],
    'most_severe_injury': [most_severe_injury],
    'crash_day_of_week': [crash_day_of_week],
    'crash_month': [crash_month],
    'damage': [damage],
    'num_units': [num_units],
    'intersection_related_i': [intersection_related_i],
    'injuries_total': [injuries_total],
    'injuries_fatal': [injuries_fatal],
    'injuries_incapacitating': [injuries_incapacitating],
    'injuries_non_incapacitating': [injuries_non_incapacitating],
    'injuries_reported_not_evident': [injuries_reported_not_evident],
    'injuries_no_indication': [injuries_no_indication],
    'road_defect': [road_defect],
    'trafficway_type': [trafficway_type]
})

# Ensure feature order matches the model's expectations
user_data = user_data[X.columns]

# Make prediction
prediction = model.predict(user_data)
print(f"\nPrediction: The predicted crash type is
{prediction[0]}.")

```



```
# Run the chatbot
```

```
accident_prediction_chatbot()
```

Welcome to the Accident Prediction Chatbot!

Please provide the following information:

Weather Condition (0=Clear, 1=Rain, 2=Snow, etc.): 0

Lighting Condition (0=Daylight, 1=Dark, etc.): 0

Road Surface Condition (0=Dry, 1=Wet, etc.): 0

Crash Hour (0-23): 3

Traffic Control Device (0=None, 1=Stop Sign, etc.): 1

First Crash Type (0=Rear-end, 1=Angle, etc.): 0

Alignment (0=Straight, 1=Curve, etc.): 1

Primary Cause (0=Speeding, 1=Distracted Driving, etc.): 1

Most Severe Injury (0=None, 1=Minor, 2=Severe, 3=Fatal): 3

Prediction: The predicted crash type is 1.